

Self-Driving Cars

Exercise 0 - Introduction

Yiyi Liao

Autonomous Vision Group
MPI-IS / University of Tübingen

October 19, 2018



University of Tübingen
MPI for Intelligent Systems

Autonomous Vision Group



Outline

- ▶ Exercises Setup
- ▶ PyTorch
- ▶ OpenAI Gym

Exercises Setup

ILIAS

- ▶ We organize the exercises using the ILIAS system
<https://ovidius.uni-tuebingen.de/ilias3>
- ▶ Exercise sheets will be available in the ILIAS system. Please be aware of the **submission deadline**.
- ▶ If you have any questions, please ask at the **Forum** on ILIAS.
- ▶ You are eligible to finish the homeworks within a group up to 2 people.
- ▶ If you don't have an university account and cannot register on ILIAS, please send us your name, email address and birthday to us yliao@tue.mpg.de, sdonne@tue.mpg.de. We will create a temporary account for you.

TCML cluster

- ▶ You are eligible to use the Training Center for Machine Learning (TCML) cluster for exercises in this lecture.
- ▶ We will create accounts for you on the cluster. Each group shares an account.
You are not supposed to apply the account by yourself.
- ▶ Please find instructions about the cluster below:
<https://docs.google.com/document/d/1AgtLy28VVZaPe79Tw0b9jjC4F1KVzffb8y1vZoURZE8/edit?usp=sharing>.

PyTorch

PyTorch

- ▶ What is PyTorch?

A Python-based scientific computing package for Deep Learning.

- ▶ Why is PyTorch?

Beginner friendly, development friendly.

- ▶ How to install?

<https://pytorch.org/get-started/locally/>

- ▶ This tutorial is for PyTorch 0.4.1

Basic Computations

Tensor

- Construct a Tensor

```
x = torch.zeros(8, 3)
print(x)
```

```
tensor([[0., 0., 0.],
        [0., 0., 0.],
        [0., 0., 0.],
        [0., 0., 0.],
        [0., 0., 0.],
        [0., 0., 0.],
        [0., 0., 0.],
        [0., 0., 0.]])
```

```
x = torch.rand(8, 3)
print(x)
```

```
tensor([[0.0984, 0.3671, 0.2543],
        [0.5016, 0.8017, 0.0918],
        [0.5081, 0.6020, 0.0867],
        [0.5313, 0.4571, 0.1624],
        [0.4231, 0.3993, 0.2713],
        [0.8978, 0.6039, 0.8519],
        [0.4829, 0.6648, 0.8295],
        [0.9060, 0.2132, 0.4110]])
```


Basic Computations

Operations

- Multiple syntaxes, e.g. Addition

```
y = torch.rand(8, 3)
```

```
print(x + y)
```

```
print(torch.add(x, y))
```

```
# providing an output tensor as argument
```

```
result = torch.empty(5, 3)
```

```
torch.add(x, y, out=result)
```

```
print(result)
```

```
# adds x to y
```

```
y.add_(x)
```

```
print(y)
```

```
tensor([[0.5029, 0.9885, 0.5179],  
        [1.1603, 0.9705, 0.3927],  
        [1.1400, 1.0469, 1.0013],  
        [1.0675, 1.0678, 0.7733],  
        [1.1638, 1.2685, 0.3746],  
        [1.7973, 1.5099, 1.6089],  
        [1.2948, 1.2600, 1.2625],  
        [1.3960, 0.9980, 0.7105]])
```

- Other operations including transposing, indexing, slicing, linear algebra etc. at <https://pytorch.org/docs/stable/torch.html>

Basic Computations

Bridge to Numpy

- PyTorch → Numpy

```
a = torch.ones(5)
b = a.numpy()
```

- Numpy → PyTorch

```
a = np.ones(5)
b = torch.from_numpy(a)
```

- Tensors can only be converted to Numpy when they are on CPU

Basic Computations

Difference to Numpy

► GPU acceleration

```
# let us run this cell only if CUDA is available  
# We will use 'torch.device' objects to move tensors in and out of GPU  
if torch.cuda.is_available():  
    device = torch.device("cuda")           # a CUDA device object  
    y = torch.ones_like(x, device=device)   # directly create on GPU  
    x = x.to(device)                        # or use 'x.to("cuda")'  
    z = x + y  
    print(z)  
    print(z.to("cpu", torch.double))        # 'x.to' can change dtype
```

```
tensor([2.9218], device='cuda:0')  
tensor([2.9218], dtype=torch.float64)
```

Basic Computations

Difference to Numpy

- ▶ GPU acceleration
- ▶ Automatic differentiation for all operations on Tensors

```
x = torch.ones(2, 2, requires_grad=True)
y = x + 2
z = y * y * 3
out = z.mean()


out.backward()

print(x.grad)
```

```
tensor([[4.5000, 4.5000],
        [4.5000, 4.5000]])
```


Neural Networks

- ▶ 1. Define a neural network with learning parameters
- ▶ 2. Process input through the network and compute loss
- ▶ 3. Propagate gradients back into the network's parameters
- ▶ 4. Update the weights of the network




Neural Networks

- ▶ Model → `torch.nn.Module`
- ▶ Loss → `torch.nn.Module.Loss`
- ▶ Optimizer → `torch.optim`
- ▶ Data → `torch.utils.data`



Neural Networks

Model



Neural Networks

Model

- ▶ Define a model as a class that inherits from `torch.nn.Module`

```
class Net(nn.Module):
```

- ▶ Define layers in the `__init__()` method

```
    def __init__(self):  
        ...
```

- ▶ Define computation flow given an input `x` in the `forward()` method

```
    def forward(self, x):  
        ...
```

- ▶ `backward()` is automatically defined

Neural Networks

Model

```
import torch
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):

    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 6, 5)
        self.conv2 = nn.Conv2d(6, 16, 5)

    def forward(self, x):
        x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        return x
```

Neural Networks

Model

- ▶ PyTorch contains a branch of standard layers which are also subclasses of `torch.nn.Module`:

- | | |
|--------------------------|---|
| ▶ Convolution layers | <code>nn.Conv2d(C_{in}, C_{out}, K)</code> |
| ▶ Pooling layers | <code>nn.MaxPool2d(K)</code> |
| ▶ Non-linear activations | <code>nn.ReLU()</code> |
| ▶ Normalization layers | <code>nn.BatchNorm2d(N)</code> |
| ▶ Linear layers | <code>nn.Linear(C_{in}, C_{out})</code> |
| ▶ ... | ... |

- ▶ It is also easy implement custom layers:

<https://pytorch.org/docs/stable/notes/extending.html#extending-torch-nn>

Neural Networks

Loss

- ▶ Loss function returns a non-negative value J measuring the distance between network estimation and the ground truth
- ▶ PyTorch contains a branch of loss functions which are also subclasses of `torch.nn.Module`:
 - ▶ L1Loss
 - ▶ MSELoss
 - ▶ CrossEntropyLoss
 - ▶ NLLLoss
 - ▶ SmoothL1Loss
 - ▶ ...

Neural Networks

Loss

- Example of using a loss function

```
loss = nn.CrossEntropyLoss()  
input = torch.randn(3, 5, requires_grad=True)  
target = torch.empty(3, dtype=torch.long).random_(5)  
output = loss(input, target)  
output.backward()
```

Neural Networks

Optimizer

- ▶ Optimizer decides how to update the parameters in the model, e.g.

$$\theta = \theta - \eta \nabla J(\theta)$$

- ▶ PyTorch implements a set of optimization algorithms in `torch.optim`:
 - ▶ SGD
 - ▶ Adam
 - ▶ LBFGS
 - ▶ ...

Neural Networks

Optimizer

- ▶ 1. Construct an Optimizer

```
optimizer = optim.SGD(model.parameters(), lr = 0.01, momentum=0.9)
```

- ▶ 2. Take an optimization steps for every batch/sample

```
for input, target in dataset:  
    # clear saved gradients before computing gradient for the new batch  
    optimizer.zero_grad()  
  
    output = model(input)  
    loss = loss_fn(output, target)  
  
    loss.backward()  
  
    # update parameters in model  
    optimizer.step()
```

Neural Networks

Data

- ▶ PyTorch provide `Dataset`, `DataLoader` in `torch.utils.data` that allows batching data, shuffling data and load data with multiple processes. Good tutorial at https://pytorch.org/tutorials/beginner/data_loading_tutorial.html
- ▶ For small scale of dataset it is fine to implement your own data loader.

Neural Networks

Saving Models

- Save/Load `state_dict` (Recommended)

```
# save
torch.save(model.state_dict(), PATH)

# load
model = TheModelClass(*args, **kwargs)
model.load_state_dict(torch.load(PATH))
model.eval()
```

- Save/Load entire model

```
# save
torch.save(model, PATH)

# load
# Model class must be defined somewhere
model = torch.load(PATH)
model.eval()
```


Neural Networks

Examples

- ▶ Toy Regression Network
- ▶ Toy Classification Network

References

- ▶ PyTorch official tutorials: <https://pytorch.org/tutorials/>
- ▶ Stanford Course on Deep Learning for Computer Vision:
http://cs231n.stanford.edu/slides/2018/cs231n_2018_lecture08.pdf
- ▶ NTU Machine Learning Course: <https://www.slideshare.net/lymanblueLin/pytorch-tutorial-for-ntu-machine-learning-course-2017>
- ▶ PyTorch tutorial with code examples:
<https://github.com/MorvanZhou/PyTorch-Tutorial>

OpenAI Gym

OpenAI Gym

- ▶ What is OpenAI Gym?

A python based toolkit for developing and comparing RL algorithms.


- ▶ Why is OpenAI Gym?

Benchmarks for RL algorithms, standardization of environments.


- ▶ How it works?

Demo


Examples




(a) Atari




(b) Box2D




(c) Controls



(d) MuJoCo




(e) Robotics



(f) Toy Texts


Reinforcement Learning

- ▶ 1. Create an Environment and initialize the State
- ▶ 2. Iteratively take an Action and observe the State, with the goal to maximize the Reward



Reinforcement Learning

- ▶ 1. Create an Environment and initialize the State
- ▶ 2. Iteratively take an Action and observe the State, with the goal to maximize the Reward



Work flow

```
import gym
env = gym.make('CarRacing-v0')
env.reset()
for _ in range(1000):

    env.render()

    # take a random action
    env.step(env.action_space.sample())
```


Work flow

Create an Environment

- ▶ Each gym environment has a unique name
- ▶ To create an environment from the name use the

```
env = gym.make(env_name)
```

- ▶ For example, to create a CarRacing environment:

```
env = gym.make('CarRacing-v0')
```

Work flow

Initialize State

- ▶ Used to reinitialize a new episode
- ▶ Returns the initial state

```
init_state = env.reset()
```

Work flow

Take an action

- ▶ Performs the specified action and returns the resulting state
- ▶ The main method your agent interacts with

```
step(action) -> (next_state ,  
                 reward ,  
                 is_terminal ,  
                 debug_info)
```

Work flow


Take an observation: Render

- ▶ Optional method
- ▶ Used to display the state of your environment
- ▶ Useful for debugging and qualitatively comparing different agent policies

```
env.render()
```

Car Racing

- Randomly generated tracks



Car Racing

- `action_space`: three continuous values, including steer, gas, brake

```
self.action_space = spaces.Box( low=np.array([-1,0,0]),  
                                high=np.array([+1,+1,+1]),  
                                dtype=np.float32)
```

- `observation_space`: color image

```
self.observation_space = spaces.Box(low=0, high=255,  
                                     shape=(STATE_H, STATE_W, 3),  
                                     dtype=np.uint8)
```

- `reward`: $R = N_{visited_tile} * \frac{1000}{N_{all_tile}} - N_{frame} * 0.1$

Car Racing



Car Racing



► Reward

Car Racing



- Reward
- Car speed

Car Racing



- ▶ Reward
- ▶ Car speed
- ▶ Wheel speed

Car Racing



- ▶ Reward
- ▶ Car speed
- ▶ Wheel speed
- ▶ Joint angle

Car Racing



- ▶ Reward
- ▶ Car speed
- ▶ Wheel speed
- ▶ Joint angle
- ▶ Angular Velocity

Car Racing

- ▶ We will create a leaderboard for the exercises
- ▶ Evaluation metrics:
 - ▶ $R = N_{visited_tile} * \frac{1000}{N_{all_tile}}$ given a fixed N_{frame}
 - ▶ $R = R - 100$ if the car get too far away from the track
- ▶ Submit your code and we will evaluate it on our server

References

- ▶ Official document: <https://gym.openai.com/docs>
- ▶ Source code: <https://github.com/openai/gym>
- ▶ https://katefvision.github.io/10703_openai_gym_recitation.pdf

Exercise 00

Install PyTorch locally

- Recommended to install with Anaconda

PyTorch Build	Stable		Preview			
Your OS	Linux		Mac		Windows	
Package	Conda		Pip		LibTorch	Source
Language	Python 2.7	Python 3.5	Python 3.6		Python 3.7	C++
CUDA	8.0		9.0		9.2	None
Run this Command:	<code>conda install pytorch torchvision -c pytorch</code>					

Exercise 00

Install OpenAI Gym locally

- ▶ Python 3.5+
- ▶ Download `sdc_gym.zip` from **ILIAS**, unzip and enter the folder
- ▶ Install the box2d package

```
pip install -e '.[box2d]'
```

- ▶ Please use the code we provided as there are some modifications compared to the official version.

Exercise 00

Install OpenAI Gym locally

- ▶ It is not recommended, but if you really want to use the official code from <https://github.com/openai/gym>, please add these two changes to `gym/envs/box2d/car_racing.py`
- ▶ Change line 336 to:

```
zoom = ZOOM*SCALE    # DO NOT Animate zoom first second
```

- ▶ Insert after line 357:

```
if "dispatch_events_called" not in self.__dict__ and mode == 'state_pixels':  
    win.dispatch_events()  
    self.dispatch_events_called = True
```

Exercise 00

Work on Cluster

- Download the Singularity image and copy it to your home directories on the cluster:

<https://owncloud.tuebingen.mpg.de/index.php/s/TNJS7Y7bXdZJfZ4>

It is an environment that contains everything you need for the exercises of this lecture such as PyTorch, OpenAI Gym.

- Run the code under the Singularity environment:

```
singularity exec ~/sdc_gym.simg python3 your_python_file.py
```

Questions?