



SELF-DRIVING CARS

EXERCISE 2 – REINFORCEMENT LEARNING

Release date: Fri, 23. November 2018 - **Deadline for Homework: Wed, 19. December 2018 - 21:00**

For this exercise you need to submit a **.zip** folder containing your report as a **.pdf** file (up to 3 pages), your best trained model **agent.pt** and your **.py** code files.

As in the previous exercise, please use the provided code templates. You may not change the `evaluate` function in `deepq.py`. Changes to the gym environment or additionally installed packages will also not be considered. Comment your code clearly and structure it well.

Note of caution: Training a reinforcement learning agent takes time. As this exercise requires you to train several agents, it may be helpful to start working on this exercise sheet reasonably early.

2.1 Base Implementation (2+2+2+2+2 Points)

The aim of this exercise is to implement a Deep Q-Learning agent for the CarRacing environment.

- a) **Deep Q-Network:** Implement a Deep Q-Network and its forward pass in the `DQN` class in `model.py`. Your network should take a single frame as input. In addition, you may again utilize the `extract_sensor_values` function. Describe your architecture in your report. Would it be a problem if we only trained our network on a crop of the game pixels, which would not include the bottom status bar and would not use the extracted sensor values as an additional input to the network? Why do we not need to use a stack of frames as input to our network as proposed to play Atari games in the original Deep Q-Learning papers [1][2]? *Hint: Think about Markov.*
- b) **Deep Q-Learning:** Implement the functions `perform_qlearning_step` and `update_target_net` in `learning.py`. The first should run a single Deep Q-Learning update step while the second should update the target network with the weights of the policy network. Why do we utilize fixed targets and sample our training data from a replay memory instead of using a batch of past consecutive frames? *Hint: Make sure to clip your gradients[1].*
- c) **Action selection:** In `action.py`, implement the functions `select_exploratory_action` and `select_greedy_action` which select an action according to an exploratory ϵ -greedy strategy or a greedy strategy (i.e. the action which maximizes the Q -values of your policy network). Why do we need to balance exploration and exploitation in a reinforcement learning agent and how does the ϵ -greedy algorithm accomplish this?
- d) **Training:** Train a Deep Q-Learning agent using the `train_racing.py` file with the provided default parameters. Describe your observations of the training process. In particular, look at the generated loss and reward curves and describe how they develop over the course of training. Some points of interests should be: How quickly is the agent able to consistently achieve positive rewards? What is the relationship between the ϵ -greedy exploration schedule and the development of the cumulative reward which the agent achieves over time? How does the loss curve compare to the loss curve that you would expect to see on a standard supervised learning problem (i.e. gradual decrease)?
- e) **Evaluation:** Evaluate the trained Deep Q-Learning agent by running the `evaluate_racing.py` script. Observe the performance of the agent by running the script on your local machine. Where does the agent do well and where does it struggle? How does its performance compare to the imitation learning agent you have trained for Exercise 1? Discuss possible reasons for the observed improvement/decline in performance compared to your imitation learning agent from Exercise 1.

Important: Make sure you have a working baseline implementation, in which the trained agent is often able to achieve positive scores as well as take some corners before moving on to work on Section 2.2.

2.2 Further Investigations and Extensions (2+2+2+2+2 Points)

Please note that all further investigations and extensions should be made w.r.t. the baseline agent which you have implemented in the previous section. Except for your best solution in part e), you should not combine subsequent extensions but investigate the effect of each aspect separately on its own.

- a) **Discount factor γ :** Investigate the influence of the discount factor γ . Why do we typically use a discount factor in reinforcement learning problems and in which cases would it be a problem not to use a discount factor (i.e. $\gamma = 1$)? What effect does an increase/decrease of γ from its default of 0.99 have on the behavior and evaluation score of your agent?
- b) **Action repeat parameter:** Describe the reasoning behind the use of an `action_repeat` parameter. Why might it be helpful to repeat each action several times? By default, this value is set to 4. What is the effect on the training progress (look at your training plots) and your evaluation performance if this parameter is increased or decreased? Discuss and interpret your findings.
- c) **Action space:** By default, the agent uses a 4-dimensional action set of left-turn, right-turn, brake and acceleration (see `get_action_set` function in `action.py`). Investigate the addition of a null action (`[0,0,0]`) as well as more fine-grained turning, braking or acceleration actions. What is the effect on the agent's driving style as well as its evaluation score? Which additional actions lead to an improvement in the agent's performance and which do not? Why might it not always be helpful to increase an agent's action space? In general, why are Deep Q-Networks limited to a discrete set of actions and what solutions exist to overcome this limitation?
- d) **Double Q-learning:** One problem with the standard Deep Q-Learning approach is an overestimation of Q -values. A proposed solution to this problem is the double Q-learning algorithm [3]. Read this paper and shortly summarize the reason for the overestimation of Q -values as well as how the double Q -learning algorithm proposes to solve this issue. Implement the double Q-learning algorithm and evaluate its effect on the training and evaluation performance of your agent. *Important: Make sure to include your double Q-learning implementation in your submitted code.*
- e) **Best solution:** Finally, putting together your previous findings as well as any more ideas you might have, fine-tune and evaluate your best-performing agent. How is this agent constructed and trained? In which aspects has its performance improved over your baseline agent from Section 2.1 and where does it still exhibit sub-optimal behavior? Briefly sketch out an idea for overcoming the main remaining issue with your agent's training or performance (*this does not need to be implemented*). Save the weights of your best performing agent as `agent.pt` and include it in your submission.

2.3 Competition (0 Points)

With each exercise sheet you are welcome to participate in a non-graded competition! Run `evaluate_racing.py` using your current model and submit your obtained evaluation score to <https://goo.gl/forms/dvtvIMEUsIAHQYOW2>. You can see the rankings here: <https://docs.google.com/spreadsheets/d/1GGd4ZhTX3OnhIq2gH4L6jzT0sOpSOIGg4bD1-PFHjG4/edit?usp=sharing>.

For a final ranking, we will run the `evaluate` function on a secret set of tracks for every submission. For each track, the cumulative reward after 600 frames is used as the performance measure. Evaluation is stopped early if the agent is `done` (e.g. has left the track). The overall score is computed by taking the mean cumulative reward from all validation tracks. The reward function is defined as -0.1 for every frame and $+1000/N$ for every track tile visited, where N is the total number of tiles in track. The winners for each of the 3 exercise sessions will be asked to present their approach in the last lecture. Good luck!

2.4 References

- [1] <https://storage.googleapis.com/deepmind-media/dqn/DQNNaturePaper.pdf>
- [2] <https://arxiv.org/pdf/1312.5602.pdf>
- [3] <https://arxiv.org/pdf/1509.06461.pdf>